

Modelling in Physical Geography

Excercise 2

Konsta Happonen

Plotting

When doing analysis, it is crucial to get to know one's data. The fastest way to do this is “eyeballing”, or visual inspection. Graphical representations of the data can be used to quickly gain information about the variation within and between variables. For this excercise we will be using one of the R default datasets from the `boot` package called `brambles`, which contains the locations and age classes of raspberry stems in a square. For more information about the dataset, type `help(boot::brambles)` on the command line.

```
### Modelling in Physical Geography
### Week 1 - Excercise 2
brambles <- boot::brambles # brambles is provided by the boot package.
str(brambles)

## 'data.frame':   823 obs. of  3 variables:
## $ x : num  0.677 0.676 0.681 0.683 0.776 0.794 0.944 0.948 0.983 0.986 ...
## $ y : num  0.001 0.022 0.031 0.038 0.028 0.033 0.011 0.01 0.077 0.084 ...
## $ age: num  0 0 0 0 0 0 0 0 0 0 ...
```


One discrete variable

Barplots

To see the distribution of one discrete variable, we can use bar plots. We can treat the `brambles` age as a discrete variable, and plot it. The (horizontal) x-axis of the resulting plot lists all the unique ages in our data, and the (vertical) y-axis represents their counts. The arguments `main`, `xlab` and `ylab` control the plot title and axis labels. Use them to make your graph more informative.

```
ages <- table(brambles$age) # Summarize the frequency of each age class
barplot(ages,
  main = "Age distribution of raspberry stems",
  xlab = "Age (years)",
  ylab = "n") # Plot them.
```

Age distribution of raspberry stems



One continuous variable


The distribution of a continuous variable is most easily visualized by histograms and boxplots.

Histograms

A histogram discretizes the continuous variable to discrete classes, and then represents their frequencies with rectangles. The area of each rectangle is directly proportional to the frequency of that class. The argument **breaks** can be used to control the coarseness of the discretisation.


```
hist(brambles$x) # plot the distribution of raspberries along the x-axis
```

Histogram of brambles\$x



```
hist(brambles$x, breaks=100)
```


Histogram of brambles\$x




Boxplots

A box plot visualizes the *quartiles* of a set of numbers. Quartiles (Q1,Q2,Q3) are calculated by first arranging the values according to their value. Q2, also called the *median*, is the middle number between the minimum and maximum values. Accordingly, Q1 is the middle number between the minimum and Q2, and Q3 is the middle number between Q2 and the maximum. The boxplot displays the minimum, Q1, Q2,Q3, and the maximum. In addition, any data points more than 1.5 interquartile ranges (distance between Q1 and Q3) away from the nearest quartile are displayed as separate points. The definition of outliers can be tweaked with the argument **range**. Boxplots are especially useful for identifying outliers (abnormally extreme values, often indicative of some kind of error in data collection) and comparing multiple variables on the same scale.

```
boxplot(brambles[,c("x","y")])
```



```
boxplot(brambles[,c("x","y")], range=0.5)
```




Two continuous variables

Often what is truly interesting in the relationship between our variables. To visualise these we can use scatterplots and lineplots

Scatterplots

In scatterplots, the values of two continuous values are mapped on the x- and y-axes. Each data point is represented a point in the plot. If the two variables are a “response” and an “explanatory” variable, the convention is to map the response on the y-axis. The argument `pch` can be used to change the point symbol, and the argument `col` to change the color of the points.


```
## Plot the spatial distribution of data with filled circles, colored by age
plot(y~x, data = brambles, col = factor(brambles$age), pch = 16)
```



Lineplots

Lineplots are good for displaying trends. The `brambles` data is ill-suited for demonstrating this, so we will create some trends. We will use the command `seq()` to create regular sequences of data. We specify `type = "l"` in the plotting command to create a line plot. We can then use `lines()` to add another line in the plot. We can also use `abline()` to create additional, straight lines.

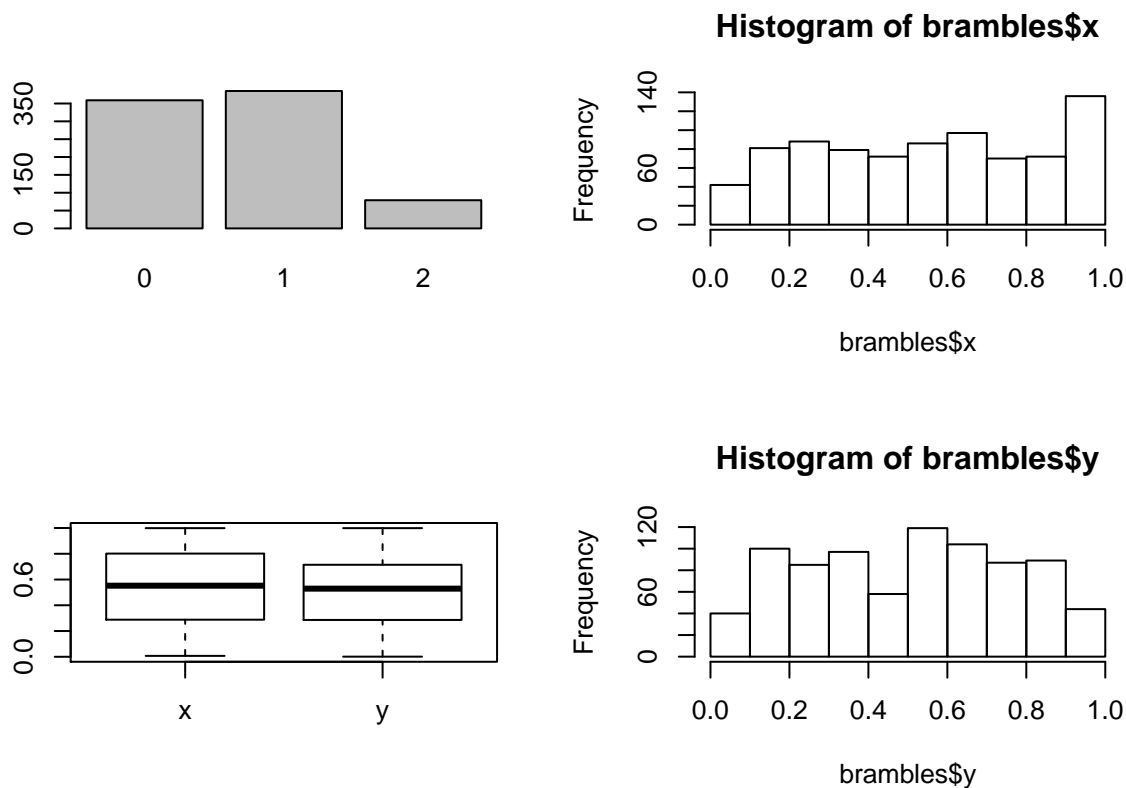
```
x <- seq(from = 1, to = 100) # equal to 1:100
y <- seq(from = 1, to = 1000, by = 10)
plot(y~x, type = "l") # create a line plot
lines(y-200~x, col = "tomato") # add a new line
abline(h = 500, lty = 2) # add a horizontal line with a new linetype
```



Multiple plots in the same window

You can modify the options of the graphics window to display multiple plots at the same time. You do this with the `par(mfrow = c(rows, columns))` command, replacing `rows` and `columns` with suitable integers. To revert your options, close the graphics window or run the command `par(mfrow = c(1, 1))`.

```
par(mfrow = c(2, 2))
barplot(ages)
hist(brambles$x)
boxplot(brambles[, c("x", "y")])
hist(brambles$y)
```



```
par(mfrow = c(1,1))
```

Saving plots

To export your plots from R, use the commands `png()` to save as raster graphics, and `pdf()` to save as vector graphics. Prefer vector graphics, as they generally take less space and retain their resolution.

The `png()` takes as its arguments the width and height of the image in pixels. The default size is quite small, so do adjust.

```
png(filename = "enigmatic_expression.png", width = 700, height = 700)
plot(x = c(2,8), y = c(7,7), xlim = c(0,10), ylim = c(0,10), col = "blue")
lines(c(6,4,3,4,6)~c(1,3,5,7,9), col = "red")
dev.off()
```

```
## pdf
## 2
```

The `pdf()` command's width and height arguments are in inches by default.

```
pdf(file = "vital_organ.pdf", width = 7, height = 7)
plot(x = c(1:9),
     y = c(7,8.5,9,8,7,8,9,8.5,7),
     xlim = c(0,10),
     col = "red",
     type = "l",
```

```
ylim = c(1,10))
lines(x = c(1,5,9), y = c(7,1,7) , col = "red")
dev.off()
```

```
## pdf
## 2
```

The images should have been saved in your working directory. Inspect them now.

Important summary statistics

We have already been introduced to means, medians, and quartiles. We will now take a closer look of **quantiles** and **standard deviations**.

Quantiles

Quantiles are the break-points when dividing a continuous-valued sample that has been ordered by its value into frequency classes. Quartiles are a special case of quantiles, as are deciles (each class comprises 10% of the samples) and percentiles. Quantiles are calculated with the function `quantile()`. The argument `probs` defines the quantiles you want as fractions (percentage / 100%). For example, `quantile(x, probs = 0.1)` gets you the value that is bigger than 10% of `x`. Quantiles are important in statistical significance testing, as we shall later see.

```
our_sample <- (1:100)^2 # generate some data
our_sample
```

```
## [1] 1 4 9 16 25 36 49 64 81 100 121
## [12] 144 169 196 225 256 289 324 361 400 441 484
## [23] 529 576 625 676 729 784 841 900 961 1024 1089
## [34] 1156 1225 1296 1369 1444 1521 1600 1681 1764 1849 1936
## [45] 2025 2116 2209 2304 2401 2500 2601 2704 2809 2916 3025
## [56] 3136 3249 3364 3481 3600 3721 3844 3969 4096 4225 4356
## [67] 4489 4624 4761 4900 5041 5184 5329 5476 5625 5776 5929
## [78] 6084 6241 6400 6561 6724 6889 7056 7225 7396 7569 7744
## [89] 7921 8100 8281 8464 8649 8836 9025 9216 9409 9604 9801
## [100] 10000
```

```
mean(our_sample)
```

```
## [1] 3383.5
```

```
median(our_sample)
```

```
## [1] 2550.5
```

```
quantile(our_sample)# the default is to calculate minimum, maximum and quartiles.
```

```
## 0% 25% 50% 75% 100%
## 1.00 663.25 2550.50 5662.75 10000.00
```




```
## this range contains 95 % of our sample.
interval <- quantile(our_sample, p = c(0.025,0.975))
interval
```

```
##      2.5%      97.5%
##    12.325 9511.375
```

```
hist(our_sample)
abline(v=mean(our_sample), col = "red") # plot the mean with red
abline(v=interval, col = "blue") # plot the quartiles in blue
```

Histogram of our_sample



Standard deviation

Standard deviation is a measure of spread for continuous variables. Sample standard deviation is calculated by dividing the sum of squared deviations from the sample mean by $N-1$ and taking the square root (`sqrt()`) of that. Or then by letting R do it for you with the function `sd()`.

```
sqrt(sum(((our_sample-mean(our_sample))^2))/(length(our_sample)-1))
```

```
## [1] 3024.356
```

```
sd(our_sample)
```

```
## [1] 3024.356
```

For a normally distributed variable, 95% of the values lie within two standard deviations of the mean.

Generating data

It is surprising how often one needs to generate data. When producing random data, they are sampled from some *distribution*. In this exercise we will cover the uniform and normal distributions.

Uniform distribution


The uniform distribution is specified with only two parameters: the **minimum** and the **maximum**. As the name implies, the probability distribution between these two points is uniform; all values have exactly equal probability of being represented in the sample. Random samples from a uniform distribution are produced with the `runif()` command.

```
## 100 draws from a uniform(0,1) distribution.
uniform <- runif(n = 100, min = 0, max = 1)

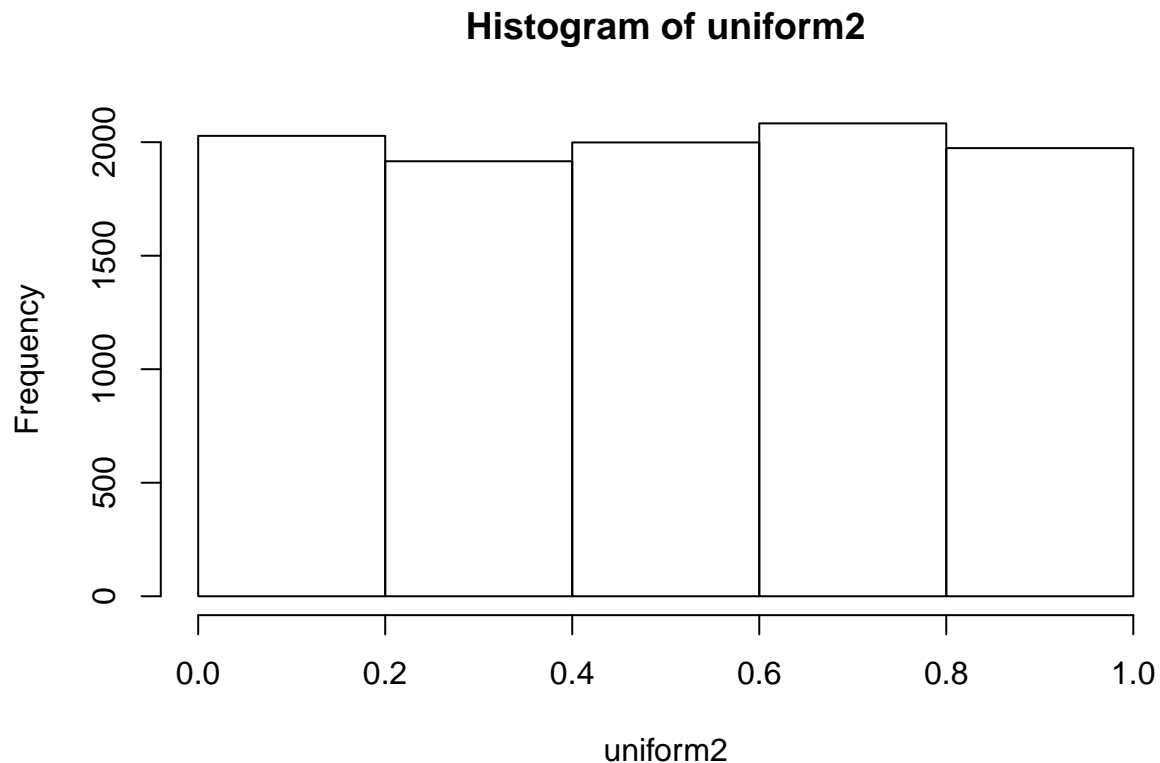
## 10000 draws from a uniform(0,1) distribution.
uniform2 <- runif(n = 10000, min = 0, max = 1)

## The sample distribution is not exactly uniform because of sampling error
hist(uniform, breaks = 5)
```

Histogram of uniform



```
hist(uniform2, breaks = 5) # much closer
```




Normal distribution

The normal distribution is one of the most important distributions there is. Its probability distribution is the famous Gauss' bell-shaped curve. The normal distribution is parametrized with a mean and standard deviation parameter. The standard deviation is a *scale parameter*, which defines the variability of sampled values. Random normal deviates are produced with the command `rnorm()`. The normal distribution is especially important because of a thing called *the central limit theorem* (CLT). The CLT states that the means of samples taken from a distribution will follow a normal distribution almost irrespective of the shape of the original distribution. CLT is an important part of statistical significance testing, as we shall later see.


```
hist(rnorm(1000, mean = 0, sd = 1), xlim = c(-10,10)) # 1000 normal deviates, sd = 1
```

Histogram of `rnorm(1000, mean = 0, sd = 1)`



```
hist(rnorm(1000, mean = 0, sd = 2), xlim = c(-10,10)) # same, but sd 2. The distribution is wider
```


Histogram of `rnorm(1000, mean = 0, sd = 2)`



Scaling variables

Comparing the distribution of two variables is hard if they are not on the same scale. In this case it can help if you *normalize* or *scale* the variables. Scaling is done by subtracting from each member of a sample the sample mean, and then dividing them by the sample standard deviation. The resulting variable will have zero mean and a standard deviation of one. Comparing variables on different scales thus becomes possible. Scaling can be done with the command `scale()`. Be aware, that because standard deviations are sensitive to outliers, scaling is as well. Scaling variables can be important in regression analysis, where it helps in the comparison and calculation of standardized effect sizes.

```
scaled_b <- scale(brambles) # Scale all variables in brambles
boxplot(scaled_b) # compare the variables on a common scale
```



Sampling

You can also randomly sample any object you have created in R with the command `sample()`. The main arguments are `x`, the sampled object, `size`, the sample size, `prob`, a vector of weights if some values are to be drawn more often than others. In addition, the argument `replace` defines whether to sample with or without replacement. Next, We will go through the differences of these two sampling methods.

Sampling without replacement

Sampling without replacement is the analogue of picking a piece of candy from a bag and eating it: you cannot eat (sample) the same candy twice. Thus when sampling without replacement, the maximum sample size is the number of candies (or observations or whatever).

```
sample(1:10, size = 5, replace = FALSE) # sample without replacement
```

```
## [1] 2 1 5 4 6
```

```
## sample(1:10, size = 11, replace = FALSE) # error: not enough samples
```

Sampling with replacement

Sampling with replacement on the other hand is comparable to picking a candy from the bag, inspecting it, putting it back again and reshuffling before taking another candy. Because the samples are replaced, there is no upper limit to the sample size. This also means that the same observation could be present in the sample more than once

```
sample(1:10, size = 5, replace = TRUE)
```

```
## [1] 7 8 2 3 8
```

```
sample(1:10, size = 11, replace = TRUE)
```


```
## [1] 5 1 2 3 8 8 6 4 10 3 7
```

Bootstrap confidence intervals

Let us look again at the spatial distribution of our raspberry stems. We want to know if there is a spatial trend in their distribution. More specifically, we want to know if the raspberries are evenly distributed from the left to the right, or if there are more raspberries on either side. For this we might investigate the mean of the x-coordinate of our raspberries. Since the possible range is 0–1, the value of the mean should be 0.5 if the raspberries are divided evenly on each side.

```
hist(brambles$x)
```

Histogram of brambles\$x



```
mean(brambles$x)
```

```
## [1] 0.5478202
```

The histogram and mean show slightly more raspberries on the right. But is this evidence of a spatial trend? If by trend we mean that the process that generates raspberries has a tendency of generating them more on the right than on the left then no. Samples from a continuous distribution will **never** have exactly the same mean as the distribution they were drawn from (or more specifically, will have the same mean with infinitely small probability). This is because of sampling error: a finite sample cannot represent the underlying distribution accurately. Small samples will, on average, differ more from the generating distribution. As sample size approaches infinity, sampling error tends to zero.


If we think our sample standard deviation accurately represents the standard deviation of the underlying distribution, and that raspberries are distributed identically and individually in our square, we can use bootstrapping to calculate the sampling distribution of the mean.

Bootstrapping is effectively just sampling with replacement, but doing it many times. For this we can use the R command `replicate()`. The argument `n` specifies how many times we want to replicate the command, and the argument `expr` specifies the command to be replicated. The argument `simplify` tells the command in what form we want the return values from our replicated command. We sample with replacement from the x-coordinates 999 times, and calculate the means of those samples. For each replicate, we set the sample size to the original sample size. We concatenate these samples to the original mean for a total of 1000 samples.

```
## Bootstrapping the mean
## default sample size for sample() is the original sample size.
bs <- replicate(n = 999, expr = mean(sample(brambles$x, replace = TRUE)))
bs <- c(bs, mean(brambles$x))
```

`bs` now contains the bootstrapped *sampling distribution of the mean*. We can now take quantiles of this distribution to produce *bootstrapped confidence intervals* for the mean x-coordinate of our raspberries. For example, if we want to calculate 95% confidence intervals, we exclude the 2.5% most extreme values from both ends.

```
confs <- quantile(bs, probs = c(0.025, 0.975))
hist(bs)
abline(v = confs, lty = 2)
```



The true mean of the distribution is, with 95% certainty, located somewhere between the dashed lines. The value 0.5 (our *null hypothesis*) is not contained in this interval. We can thus conclude (with 95% certainty) that raspberries have a tendency of being on the right.